TASK 2 | UNDERSTANDING ARTIFICIAL INTELLIGENCE

How Machines Learn

Exploring the Logic Behind Machine Learning Through Everyday Examples

EMAIL SPAM FILTERS

VOICE ASSISTANTS

SHOPPING RECOMMENDATIONS

PREPARED BY
Santanu Mondal

PROGRAMME
B.Tech CSE (AI & ML), Adamas University

SUBMITTED FOR
InAmigos Foundation — Task 2

FOCUS AREA
Machine Learning Fundamentals

1. Introduction

Machine Learning (ML) is the branch of Artificial Intelligence that allows computers to improve at a task by studying examples, instead of being given a fixed set of rules by a programmer. Rather than writing an explicit instruction for every possible situation, we show the system many past examples — called **training data** — and let it discover the underlying pattern on its own. Once this pattern is captured inside a **model**, the system can apply it to brand-new situations it has never seen before, producing a prediction, a classification, or a decision.

This report looks at three everyday applications of Machine Learning — **e-mail spam detection**, **voice assistants**, and **online shopping recommendations** — and explains, in simple terms, what each system learns from and how it uses that learning to help us in daily life.

2. The General Idea: A Learning Loop

Before looking at individual examples, it helps to see the common pattern that every Machine Learning system follows. Almost all ML applications — no matter how different they look on the surface — move through the same five-stage loop:

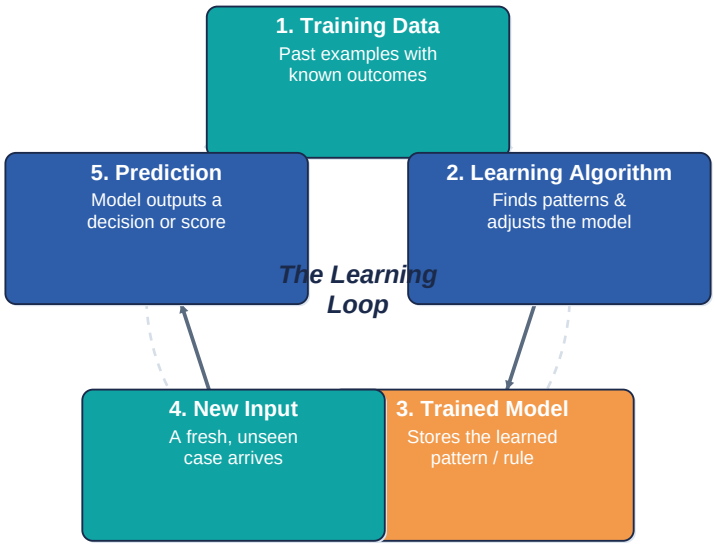

Figure 1: The general Machine Learning loop — data is used to train a model, the model makes predictions on new inputs, and feedback from those predictions is fed back in to keep improving the model over time.

01 E-mail Spam Detection

Separating genuine messages from unwanted junk mail

Every e-mail inbox uses a spam filter running quietly in the background. Its job is to look at an incoming message and decide, in a fraction of a second, whether it belongs in the Inbox or the Spam folder.

WHAT IT LEARNS FROM

Labelled example e-mails

The system is trained on millions of past e-mails that humans have already marked as **“spam”** or **“not spam.”** From these, it studies patterns such as certain keywords (e.g. “free,” “winner,” “click now”), suspicious links, sender reputation, excessive capital letters, and formatting tricks that spam messages tend to share.

HOW IT USES THAT LEARNING

Scoring every new e-mail

When a new e-mail arrives, the model checks it against the patterns it has learned and calculates a **spam probability score**. If the score crosses a set threshold, the mail is automatically moved to Spam; otherwise it lands in the Inbox. Every time a user manually marks a mail as spam or “not spam,” that correction becomes new training data.

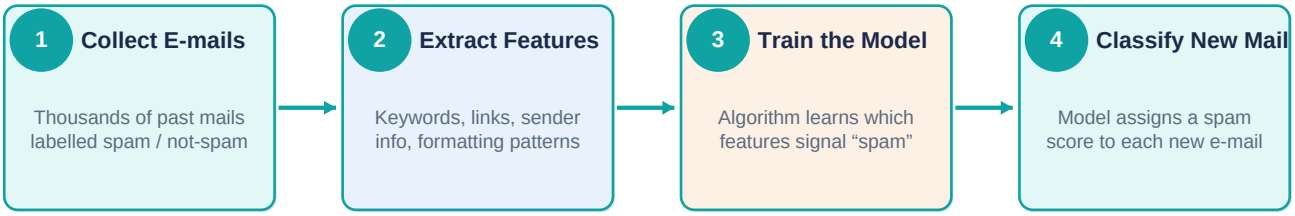

Figure 2: How a spam filter turns labelled e-mail examples into an automatic “Inbox vs. Spam” decision for every new message.

02 Voice Assistants

Turning spoken words into understood commands

Voice assistants such as Google Assistant, Siri, and Alexa listen to a spoken request, convert it into text, understand its meaning, and then carry out an action — all within a second or two. This is possible because several ML models work together behind the scenes.

WHAT IT LEARNS FROM

Huge libraries of recorded speech

The assistant is trained on massive datasets of recorded human speech paired with its correct text transcription, across different accents, speeds, and background noise levels. It also learns from labelled examples of **intent** — e.g. phrases that mean “set an alarm” versus phrases that mean “play music.”

HOW IT USES THAT LEARNING

Listen, understand, respond

When you speak, the model converts your voice into text (speech recognition), works out what you want (intent recognition), and triggers the matching action — playing a song, setting a timer, or answering a question. Each interaction, especially corrections, helps refine the model further.

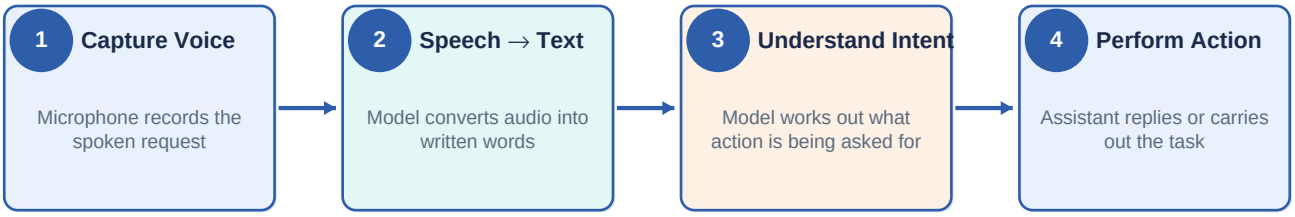

Figure 3: How a voice assistant moves from raw sound to a completed action, one learned step at a time.

03

Online Shopping Recommendations

Suggesting products you are likely to want next

Whenever an e-commerce site shows a “Customers who bought this also bought…” or “Recommended for you” section, a recommendation engine is at work, trying to predict which products best match a particular shopper.

WHAT IT LEARNS FROM

Browsing and purchase behaviour

The system studies each user’s past clicks, searches, items viewed, items purchased, ratings given, and time spent on a product page. It also compares this behaviour against **millions of other shoppers** with similar tastes, to spot patterns such as “customers who bought a phone also bought a screen protector.”

HOW IT USES THAT LEARNING

Ranking products for you

Using these learned patterns, the model ranks the entire product catalogue by how likely each item is to interest the current shopper, and displays the top matches as personalised recommendations. Every new click or purchase feeds back into the model, keeping the suggestions fresh and relevant.

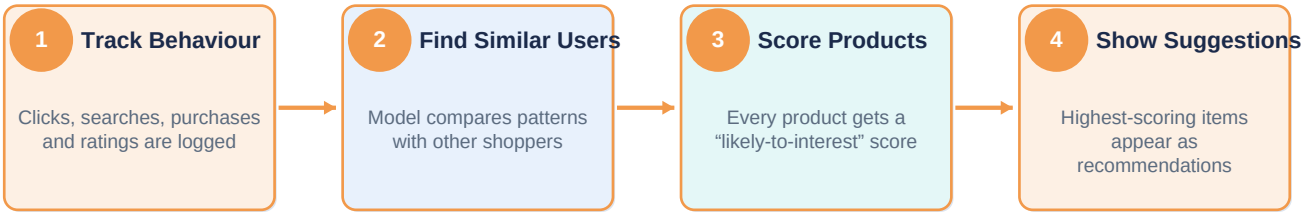

Figure 4: How a recommendation engine turns shopping behaviour into a personalised list of suggested products.

4. Comparing the Three Systems

The table below summarises what each system learns from, what kind of output it produces, and the type of learning problem it represents.

System	Learns From	Produces	Type of Task
Spam Detection	Labelled e-mails (spam / not spam)	Inbox vs. Spam decision	Classification
Voice Assistant	Recorded speech + labelled intents	Text transcript & completed action	Speech recognition + Classification
Shopping Recommendations	Browsing / purchase history of users	Ranked list of suggested products	Recommendation / Ranking

5. Key Takeaways

- **Learning from examples, not rules:** In every case, the system was never explicitly told “if the subject line contains X, mark as spam.” It figured out the pattern itself by studying many labelled examples.
- **Data quality drives accuracy:** The more varied and correctly labelled the training data (e-mails, voice recordings, or shopping histories), the more reliable the predictions become.
- **Feedback keeps the model improving:** User actions — marking a mail as spam, correcting a misheard command, or clicking a recommended product — are fed back into the system, allowing it to continuously get better over time.
- **Same loop, different domain:** Although spam filters, voice assistants, and shopping engines look completely different to a user, all three follow the same underlying cycle — collect data, train a model, predict, and improve from feedback.

6. Conclusion

Machine Learning quietly powers many tools used every day, from keeping inboxes clean to understanding spoken commands and suggesting the next product to buy. In each case, the core idea remains the same: a system is shown many real examples, it learns the pattern connecting those examples to the correct outcome, and it then applies that learned pattern to make fast, automated decisions about new, unseen situations. Understanding this simple loop — *data, learning, prediction, feedback* — is the foundation for understanding almost every Machine Learning application in use today.